

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

Data preparation

```
In [3]: df=pd.read_csv(r"C:\Users\Sonu\OneDrive\Desktop\NIT\29th- REGRESSION PROJECT\29th-
```

Explore the data


```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      18249 non-null  int64
1   Date            18249 non-null  object
2   AveragePrice    18249 non-null  float64
3   Total Volume    18249 non-null  float64
4   4046            18249 non-null  float64
5   4225            18249 non-null  float64
6   4770            18249 non-null  float64
7   Total Bags      18249 non-null  float64
8   Small Bags      18249 non-null  float64
9   Large Bags      18249 non-null  float64
10  XLarge Bags     18249 non-null  float64
11  type            18249 non-null  object
12  year            18249 non-null  int64
13  region          18249 non-null  object
dtypes: float64(9), int64(2), object(3)
memory usage: 1.9+ MB
```

```
In [5]: df.head()
```

Out[5]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042.21
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5677.40
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986.26



Missing value checking

In [6]: `df.isnull().sum()`

Out[6]:

Unnamed: 0	0
Date	0
AveragePrice	0
Total Volume	0
4046	0
4225	0
4770	0
Total Bags	0
Small Bags	0
Large Bags	0
XLarge Bags	0
type	0
year	0
region	0

dtype: int64

Dropping unnecessary columns

In [7]: `df = df.drop(['Unnamed: 0', '4046', '4225', '4770', 'Date'], axis=1)`

In [8]: `df.head()`

Out[8]:

	AveragePrice	Total Volume	Total Bags	Small Bags	Large Bags	XLarge Bags	type	year	region
0	1.33	64236.62	8696.87	8603.62	93.25	0.0	conventional	2015	Albany
1	1.35	54876.98	9505.56	9408.07	97.49	0.0	conventional	2015	Albany
2	0.93	118220.22	8145.35	8042.21	103.14	0.0	conventional	2015	Albany
3	1.08	78992.15	5811.16	5677.40	133.76	0.0	conventional	2015	Albany
4	1.28	51039.60	6183.95	5986.26	197.69	0.0	conventional	2015	Albany

Answering questions

```
In [9]: def get_avarage(df,column):
        """
        Description: This function to return the average value of the column

        Arguments:
            df: the DataFrame.
            column: the selected column.
        Returns:
            column's average
        """
        return sum(df[column])/len(df)
```

```
In [10]: def get_avarge_between_two_columns(df,column1,column2):
        """
        Description: This function calculate the average between two columns in the dat

        Arguments:
            df: the DataFrame.
            column1:the first column.
            column2:the scond column.
        Returns:
            Sorted data for relation between column1 and column2
        """

        List=list(df[column1].unique())
        average=[]

        for i in List:
            x=df[df[column1]==i]
            column1_average= get_avarage(x,column2)
            average.append(column1_average)

        df_column1_column2=pd.DataFrame({'column1':List,'column2':average})
        column1_column2_sorted_index=df_column1_column2.column2.sort_values(ascending=F
        column1_column2_sorted_data=df_column1_column2.reindex(column1_column2_sorted_i

        return column1_column2_sorted_data
```

```
In [11]: def plot(data,xlabel,ylabel):
        """
```

Description: This function to draw a barplot

Arguments:

data: the DataFrame.

xlabel: the label of the first column.

ylabel: the label of the second column.

Returns:


None

"""

```
plt.figure(figsize=(15,5))
ax=sns.barplot(x=data.column1,y=data.column2,palette='rocket')
plt.xticks(rotation=90)
plt.xlabel(xlabel)
plt.ylabel(ylabel)
plt.title(('Avarage '+ylabel+' of Avocado According to '+xlabel));
```

Which regression are the lowest and highest prices of Avocado?

```
In [12]: data1 = get_avarge_between_two_columns(df,'region','AveragePrice')
plot(data1,'Region','Price ($)')
```



```
In [13]: print(data1['column1'].iloc[-1], " is the region producing avocado with the lowest price.")
```


Houston is the region producing avocado with the lowest price.

What is the highest region of avocado production?

```
In [14]: data2 = get_avarge_between_two_columns(df,'region','Total Volume')
sns.boxplot(x=data2.column2).set_title("Figure: Boxplot repersenting outlier column")
```

```
Out[14]: Text(0.5, 1.0, 'Figure: Boxplot repersenting outlier columns.')
```

Figure: Boxplot repersenting outlier columns.




```
In [15]: outlier_region = data2[data2.column2>10000000]
print(outlier_region['column1'].iloc[-1], "is outlier value")
```

TotalUS is outlier value

Remove the outlier values


```
In [16]: outlier_region.index
data2 = data2.drop(outlier_region.index,axis=0)
```

```
In [17]: plot(data2,'Region','Volume')
```



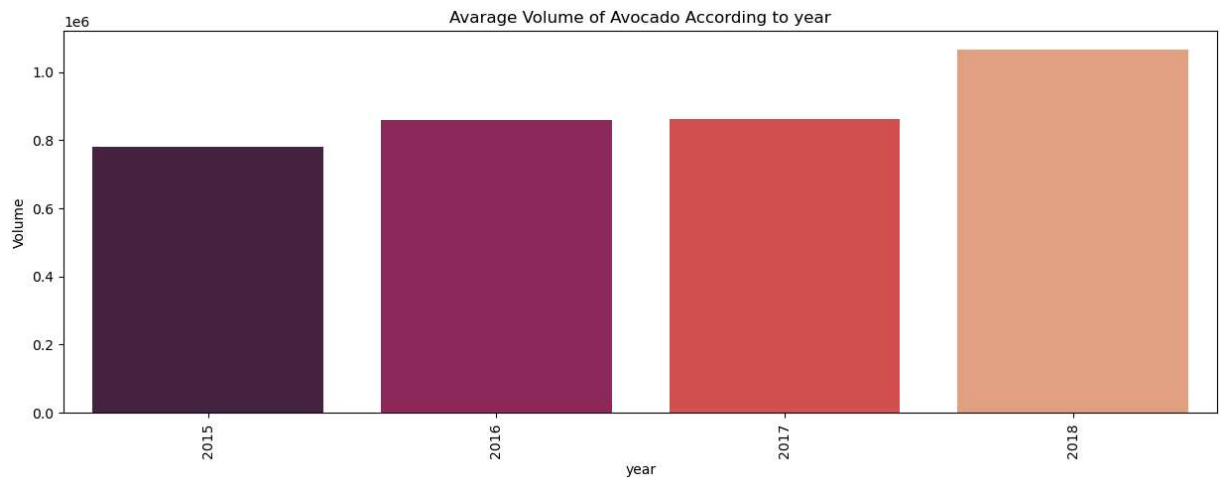
What is the average avocado prices in each year?

```
In [18]: data3 = get_avarge_between_two_columns(df, 'year', 'AveragePrice')
plot(data3, 'year', 'Price')
```



What is the average avocado volume in each year?

```
In [19]: data4 = get_avarge_between_two_columns(df, 'year', 'Total Volume')
plot(data4, 'year', 'Volume')
```



Data Modeling

```
In [20]: df['region'] = df['region'].astype('category')
df['region'] = df['region'].cat.codes

df['type'] = df['type'].astype('category')
df['type'] = df['type'].cat.codes
```

```
In [21]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   AveragePrice    18249 non-null  float64
1   Total Volume    18249 non-null  float64
2   Total Bags      18249 non-null  float64
3   Small Bags      18249 non-null  float64
4   Large Bags      18249 non-null  float64
5   XLarge Bags     18249 non-null  float64
6   type            18249 non-null  int8
7   year            18249 non-null  int64
8   region          18249 non-null  int8
dtypes: float64(6), int64(1), int8(2)
memory usage: 1.0 MB
```

```
In [22]: df.head()
```

Out[22]:

	AveragePrice	Total Volume	Total Bags	Small Bags	Large Bags	XLarge Bags	type	year	region
0	1.33	64236.62	8696.87	8603.62	93.25	0.0	0	2015	0
1	1.35	54876.98	9505.56	9408.07	97.49	0.0	0	2015	0
2	0.93	118220.22	8145.35	8042.21	103.14	0.0	0	2015	0
3	1.08	78992.15	5811.16	5677.40	133.76	0.0	0	2015	0
4	1.28	51039.60	6183.95	5986.26	197.69	0.0	0	2015	0

```
In [23]: # split data into X and y
X = df.drop(['AveragePrice'],axis=1)
y = df['AveragePrice']

# split data into training and testing dataset
X_train, X_test, y_train, y_test = train_test_split(X,
                                                    y,
                                                    test_size=0.3,
                                                    random_state=15)
```

```
In [24]: print("training set:",X_train.shape,' - ',y_train.shape[0],' samples')
print("testing set:",X_test.shape,' - ',y_test.shape[0],' samples')
```

```
training set: (12774, 8) - 12774 samples
testing set: (5475, 8) - 5475 samples
```

```
In [27]: model=LinearRegression(normalize=True)
model.fit(X_train,y_train)
```

```
-----
TypeError                                Traceback (most recent call last)
Cell In[27], line 1
----> 1 model=LinearRegression(normalize=True)
      2 model.fit(X_train,y_train)

TypeError: LinearRegression.__init__() got an unexpected keyword argument 'normalize'
```

Evaluate the results

```
In [29]: # prediction and calculate the accuracy for the testing dataset
test_pre = model.predict(X_test)
test_score = r2_score(y_test,test_pre)
print("The accuracy of testing dataset ",test_score*100)
```

```
The accuracy of testing dataset 38.58074176452857
```

```
In [30]: # prediction and calculate the accuracy for the testing dataset
train_pre = model.predict(X_train)
train_score = r2_score(y_train,train_pre)
print("The accuracy of training dataset ",train_score*100)
```

```
The accuracy of training dataset 39.7068604241106
```


In []: